

SequoiaDB Spark Yarn 部署及案例演示

1、背景

由于 MRv1 在扩展性、可靠性、资源利用率和多框架等方面存在明显的不足，在 Hadoop MRv2 中引入了资源管理和调度系统 YARN。YARN 是 Hadoop MRv2 计算机框架中构建的一个独立的、通用的资源管理系统，可为上层应用提供统一的资源管理和调度，它的引入为集群在利用率、资源统一管理和数据共享等方面带来了巨大好处。主要体现在以下几个方面：

(1) 资源利用率大大提高。一种计算框架一个集群，往往会由于应用程序数量和资源需求的不均衡性，使得在某段时间有些计算框架集群资源紧张，而另外一些集群资源空闲。共享集群模式则通过多种框架共享资源，使得集群中的资源得到更加充分的利用；

(2) 运维成本大大降低。共享集群模式使得少数管理员就可以完成多个框架的统一管理；

(3) 共享集群的模式也让多种框架共享数据和硬件资源更为方便。

2、产品介绍

巨杉数据库 SequoiaDB 是一款分布式非关系型文档数据库，可以被用来存取海量非关系型的数据，其底层主要基于分布式，高可用，高性能与动态数据类型设计，它兼顾了关系型数据库中众多的优秀设计：如索引、动态查询和更新等，同时以文档记录为基础更好地处理了动态灵活的数据类型。并且为了用户能够使用常见的分布式计算框架，SequoiaDB 可以和常见分布式计算框架如 Spark、Hadoop、HBase 进行整合。本文主要讲解 SequoiaDB 与 Spark、YARN 的整合以及通过一个案例来演示 MapReduce 分析存储在 SequoiaDB 中的业务数据。

3、环境搭建

3.1、服务器分布

服务器	服务名称
192.168.1.46	
192.168.1.47	
192.168.1.48	NameNode、DataNode、

3.2、 软件配置

操作系统: RedHat6.5

JDK 版本: 1.7.0_80 64 位

Scala 版本:

Hadoop 版本: 2.7.2

Spark 版本: 2.0

SequoiaDB 版本: 2.0

3.3、 安装步骤

1、JDK 安装

```
tar -xvf jdk-7u45-linux-x64.tar.gz -C /usr/local
cd /usr/local
ln -s jdk1.7.0_45 jdk
```

配置环境变量

```
vim ~/.bash_profile
export JAVA_HOME=/usr/local/jdk
export CLASS_PATH=$JAVA_HOME/lib:$JAVA_HOME/jre/lib
export PATH=$PATH:$JAVA_HOME/bin
source /etc/profile
```

2、Scala 安装

```
tar -xvf scala-2.11.8.tgz -C /usr/local
cd /usr/local
ln -s scala-2.11.8 scala
```

配置环境变量

```
vim ~/.bash_profile
export SCALA_HOME=/usr/local/scala
export PATH=$PATH:$SCALA_HOME/bin
```

3、修改主机 hosts 文件配置

在每台主机上修改 host 文件

```
vim /etc/hosts
192.168.1.46 node01
192.168.1.47 node02
192.168.1.48 master
```

4、 SSH 免密钥登录

在 master 节点中执行 ssh-keygen 按回车键

```
cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
```

将 master 节点中的授权文件 authorized_keys 传输到 slave 节点中

```
scp ~/.ssh/id_rsa.pub root@master:~/.ssh/
```

在 slave 节点中执行

```
cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
```

在 slave 节点中验证 SSH 免密钥登录

```
ssh master
```

5、Hadoop 集群安装

拷贝 hadoop 文件 hadoop-2.7.2.tar.gz 到/opt 目录中

解压 hadoop 安装包

```
tar -xvf hadoop-2.7.2.tar.gz
```

```
mv hadoop-2.7.2 /opt/cloud/hadoop
```

创建 hadoop 数据存储及临时目录

```
mkdir -p /opt/hadoop/data
```

```
mkdir -p /opt/hadoop/tmp
```

配置 Hadoop jdk 环境变量

```
vim hadoop-env.sh
```

```
export JAVA_HOME=/usr/local/jdk
```

编辑 core.xml 文件

```
<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://master:9000</value>
  </property>
  <property>
    <name>hadoop.tmp.dir</name>
    <value>/opt/data/tmp</value>
  </property>
  <property>
    <name>io.file.buffer.size</name>
    <value>4096</value>
  </property>
</configuration>
```

编辑 mapred-site.xml

```
<configuration>
```

```
<property>
  <name>mapreduce.framework.name</name>
  <value>yarn</value>
</property>
<property>
  <name>mapreduce.jobtracker.http.address</name>
  <value> master:50030</value>
</property>
<property>
  <name>mapreduce.jobhistory.address</name>
  <value> master:10020</value>
</property>
<property>
  <name>mapreduce.jobhistory.webapp.address</name>
  <value>master:19888</value>
</property>
</configuration>
```

编辑 hdfs-site.xml

```
<configuration>
  <property>
    <name>dfs.nameservices</name>
    <value>master</value>
  </property>
  <property>
    <name>dfs.namenode.secondary.http-address</name>
    <value> master:50090</value>
  </property>
  <property>
    <name>dfs.namenode.name.dir</name>
    <value>file:///opt/hadoop/data/name</value>
  </property>
  <property>
    <name>dfs.datanode.data.dir</name>
    <value>file:///opt/hadoop/data</value>
  </property>
  <property>
    <name>dfs.replication</name>
    <value>3</value>
  </property>
  <property>
    <name>dfs.webhdfs.enabled</name>
    <value>true</value>
  </property>
</configuration>
```

```
    </property>
</configuration>
```

编辑 yarn-site.xml

```
<configuration>
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
  <property>
    <name>yarn.resourcemanager.address</name>
    <value> master:8032</value>
  </property>
  <property>
    <name>yarn.resourcemanager.scheduler.address</name>
    <value> master:8030</value>
  </property>
  <property>
    <name>yarn.resourcemanager.resource-tracker.address</name>
    <value> master:8031</value>
  </property>
  <property>
    <name>yarn.resourcemanager.admin.address</name>
    <value> master:8033</value>
  </property>
  <property>
    <name>yarn.resourcemanager.webapp.address</name>
    <value>master:8088</value>
  </property>
  <property>
    <name>yarn.nodemanager.resource.memory-mb</name>
    <value>12288</value>
  </property>
  <property>
    <name>yarn.nodemanager.log-dirs</name>
    <value>/opt/hadoop/tmp/userlogs</value>
  </property>
</configuration>
```

启动 Hadoop

首次启动集群时，做如下操作

进入到/opt/cloud/hadoop/bin 目录中执行./hdfs namenode -format 格式化

hdfs 文件系统

进入到/opt/cloud/hadoop/sbin 目录中执行./start-all.sh 启动 hadoop 集群

6、安装 Spark 集群

拷贝 Spark 安装包到/opt 目录中，解压

```
tar -xvf spark-2.0.0-bin-hadoop2.7.tgz
```

```
mv spark-2.0.0-bin-hadoop2.7 /opt/cloud/spark
```

编辑 spark-env.sh

```
vim spark-env.sh
```

```
JAVA_HOME="/usr/jdk1.7"
```

```
SPARK_DRIVER_MEMORY="1g"
```

```
SPARK_EXECUTOR_CORES=1
```

```
SPARK_EXECUTOR_MEMORY="512m"
```

```
SPARK_MASTER_PORT="7077"
```

```
SPARK_MASTER_WEBUI_PORT="8070"
```

```
SPARK_CLASSPATH="/opt/cloud/spark/jars/sequoiadb.jar:/opt/cloud/s  
park/jars/spark-sequoiadb_2.11-2.6.0.jar"
```

```
SPARK_MASTER_IP="node03"
```

```
SPARK_WORKER_MEMORY="712m"
```

```
SPARK_WORKER_CORES=1
```

```
SPARK_WORKER_INSTANCES=1
```

```
SPARK_WORKER_DIR="/opt/data/spark/work"
```

```
SPARK_LOCAL_DIRS="/opt/data/spark/tmp"
```

```
HADOOP_HOME="/opt/cloud/hadoop"
```

```
HADOOP_CONF_DIR=$HADOOP_HOME/etc/hadoop
```

编辑 slaves

```
node02
```

```
node03
```

启动 spark 集群

进入到目录/opt/cloud/spark/sbin 目录中

```
./start-all.sh
```

Spark 成功启动后截图如下：

Spark Master at spark://node03:7077

URL: spark://node03:7077
 REST URL: spark://node03:6066 (cluster mode)
 Alive Workers: 2
 Cores in use: 2 Total, 0 Used
 Memory in use: 1424.0 MB Total, 0.0 B Used
 Applications: 0 Running, 0 Completed
 Drivers: 0 Running, 0 Completed
 Status: ALIVE

Workers

Worker Id	Address	State	Cores	Memory
worker-20161116001004-192.168.1.48-59316	192.168.1.48:59316	ALIVE	1 (0 Used)	712.0 MB (0.0 B Used)
worker-20161118215804-192.168.1.47-40327	192.168.1.47:40327	ALIVE	1 (0 Used)	712.0 MB (0.0 B Used)

Running Applications

Application ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
----------------	------	-------	-----------------	----------------	------	-------	----------

Completed Applications

Application ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
----------------	------	-------	-----------------	----------------	------	-------	----------

- 7、Spark Yarn 连接 SequoiaDB
 在 SequoiaDB 中创建集合空间、集合
 db.createCS('poc');
 db.poc.createCL('test');

```
help() for help, Ctrl+c or quit to exit
> db = new Sdb();
localhost:11810
Takes 0.6467s.
> db.createCS('poc');
localhost:11810.poc
Takes 0.26202s.
> db.poc.createCL('test');
localhost:11810.poc.test
Takes 19.459596s.
> db.poc.test.insert({id:1,name:'测试1'});
Takes 0.35498s.
> db.poc.test.insert({id:2,name:'测试2'});
Takes 0.17728s.
> db.poc.test.find();
{
  "_id": {
    "$oid": "582f0b47041f736581000000"
  },
  "id": 1,
  "name": "测试1"
}
{
  "_id": {
    "$oid": "582f0b4d041f736581000001"
  },
  "id": 2,
  "name": "测试2"
}
Return 2 row(s).
Takes 0.42897s.
```

进入到 spark 安装目录 bin 中，执行 ./spark-sql -master yarn 启动 spark sql 交互界面

```
ups with view permissions: Set(); users with modify permissions: Set(sdbadmin); groups with modify permissions: Set()
16/11/14 10:37:32 INFO Client: Submitting application application_1479090455075_0002 to ResourceManager
16/11/14 10:37:32 INFO YarnClientImpl: Submitted application application_1479090455075_0002
16/11/14 10:37:32 INFO SchedulerExtensionServices: Starting Yarn extension services with app application_1479090455075_0002 and attemptId None
16/11/14 10:37:33 INFO Client: Application report for application_1479090455075_0002 (state: ACCEPTED)
16/11/14 10:37:33 INFO Client:
  client token: N/A
  diagnostics: N/A
  ApplicationMaster host: N/A
  ApplicationMaster RPC port: -1
  queue: default
  start time: 1479091052125
  final status: UNDEFINED
  tracking URL: http://sequoia1:8088/proxy/application_1479090455075_0002/
  user: sdbadmin
16/11/14 10:37:34 INFO Client: Application report for application_1479090455075_0002 (state: ACCEPTED)
16/11/14 10:37:35 INFO Client: Application report for application_1479090455075_0002 (state: ACCEPTED)
16/11/14 10:37:36 INFO Client: Application report for application_1479090455075_0002 (state: ACCEPTED)
16/11/14 10:37:37 INFO Client: Application report for application_1479090455075_0002 (state: ACCEPTED)
16/11/14 10:37:38 INFO Client: Application report for application_1479090455075_0002 (state: ACCEPTED)
```

创建表，映射到上述 poc 集合空间中 test 集合
CREATE TABLE `test` (`id` INT, `name` STRING)

USING com.sequoiadb.spark


```
OPTIONS (
  `collection` 'test',
  `host` 'node02:11810,node03:11810',
  `serialization.format` '1',
  `collectionspace` 'poc'
);
```

查询表 test 数据，执行：

Select * from test;

```
which has no missing parents
16/11/18 22:09:15 INFO memory.MemoryStore: Block broadcast_1 stored as values in memory (estimated size 7.8 KB, fre
16/11/18 22:09:15 INFO memory.MemoryStore: Block broadcast_1_piece0 stored as bytes in memory (estimated size 4.2 K
16/11/18 22:09:15 INFO storage.BlockManagerInfo: Added broadcast_1_piece0 in memory on 192.168.1.48:42710 (size: 4.
GB)
16/11/18 22:09:15 INFO spark.SparkContext: Created broadcast 1 from broadcast at DAGScheduler.scala:1012
16/11/18 22:09:15 INFO scheduler.DAGScheduler: Submitting 1 missing tasks from ResultStage 1 (MapPartitionsRDD[7] a
CliDriver.java:376)
16/11/18 22:09:15 INFO cluster.YarnScheduler: Adding task set 1.0 with 1 tasks
16/11/18 22:09:15 INFO scheduler.TaskSetManager: Starting task 0.0 in stage 1.0 (TID 1, node02, partition 0, RACK_L
16/11/18 22:09:15 INFO cluster.YarnSchedulerBackend$YarnDriverEndpoint: Launching task 1 on executor id: 1 hostname
16/11/18 22:09:15 INFO storage.BlockManagerInfo: Added broadcast_1_piece0 in memory on node02:56899 (size: 4.2 KB,
16/11/18 22:09:15 INFO scheduler.TaskSetManager: Finished task 0.0 in stage 1.0 (TID 1) in 747 ms on node02 (1/1)
16/11/18 22:09:15 INFO cluster.YarnScheduler: Removed TaskSet 1.0, whose tasks have all completed, from pool
16/11/18 22:09:15 INFO scheduler.DAGScheduler: ResultStage 1 (processCmd at CliDriver.java:376) finished in 0.749 s
16/11/18 22:09:15 INFO scheduler.DAGScheduler: Job 1 finished: processCmd at CliDriver.java:376, took 0.814116 s
1
测试1
2
测试2
Time taken: 5.627 seconds, Fetched 2 row(s)
16/11/18 22:09:15 INFO CliDriver: Time taken: 5.627 seconds, Fetched 2 row(s)
```

进入到 yarn 管理页面查看 spark 任务



RUNNING Applications

Cluster

About

Nodes

Node Labels

Applications

NEW

NEW SAVING

SUBMITTED

ACCEPTED

RUNNING

FINISHED

FAILED

KILLED

Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved	Active Nodes	Decommissioned Nodes	Lost Nodes
4	0	1	3	2	2 GB	16 GB	0 B	2	16	0	2	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation
Capacity Scheduler	[MEMORY]	<memory:1024, vCores:1>	<memory:8192, vCores:8>

Show 20 entries

ID	User	Name	Application Type	Queue	StartTime	FinishTime	State	FinalStatus	Progress	Tracking
application_1479225968693_0004	root	SparkSQL::192.168.1.48	SPARK	default	Fri Nov 18 19:56:50 +0800 2016	N/A	RUNNING	UNDEFINED		Applicati

Showing 1 to 1 of 1 entries

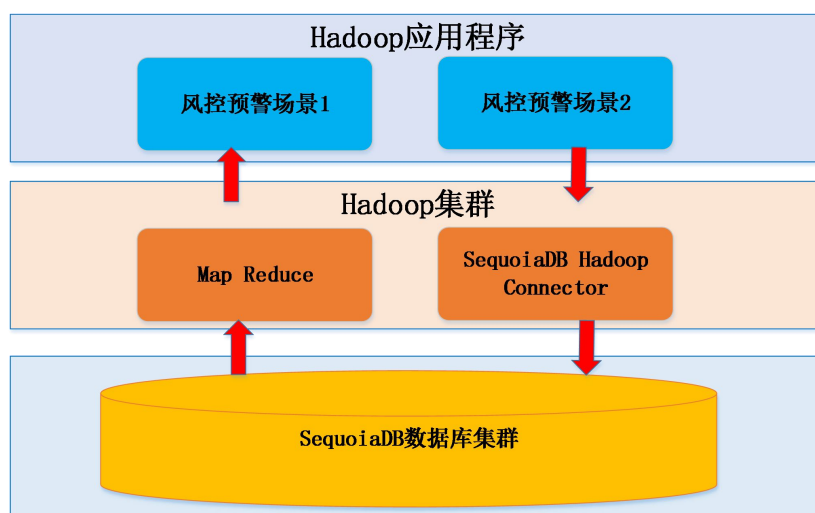
5、案例演示

为了配合司法部门的执法和银行内部的风险监管，部分商业银行对于存取款业务定制了相关预警方案，本案例以个人存取款业务高频交易来讲述 MapReduce 如何分析 SequoiaDB 中的个人交易明细数据。

具体场景为：分析同一实体柜员办理，1 小时内同一账户连续 3 笔以上支取类金额的交易账户及明细。

本演示案例采用 Hadoop Map Reduce 实现，开发语言为 Java 语言。整个测试程序分为两个部分 Map 算法和 Reduce 算法。演示程序中 Map 算法负责将同一个账号的所有对应交易明细归并在一起并输出给 Reduce 端，Reduce 端根据 Map 算法的结果运算具体的业务场景，最后将运算结果写入到 SequoiaDB 中。

具体架构如下：



Reduce 端具体算法流程如下：


```

        context.write(new Text(acct_no), value);
    }
}

```

Reduce 端算法代码如下:

```

static class TReducer extends Reducer<Text,BSONWritable,NullWritable,NullWritable>{
    private static String pattern = "yyyy-MM-dd HH:mm:ss";
    private DateFormat df = new SimpleDateFormat(pattern);

    private static int tradeNum1 = 3;
    private static int tradeTime1 = 3600;

    private static int tradeNum2 = 2;
    private static int tradeTime2 = 1800;
    private static int tradeAll = 100000;

    private Sequoiadb sdb = null;
    private CollectionSpace cs = null;
    private DBCollection cl_1 = null;
    private DBCollection cl_2 = null;
    private static String CS_NAME="";
    private static String CL_NAME_1="";
    private static String CL_NAME_2="";

    public TReducer() {
        if (null == sdb) {
            sdb = ConnectionPool.getInstance().getConnection();
        }

        if (sdb.isCollectionSpaceExist(CS_NAME)) {
            cs = sdb.getCollectionSpace(CS_NAME);
        } else {
            throw new BaseException("集合空间" + CS_NAME + "不存在!");
        }

        if (null == cs) {
            throw new BaseException("集合空间不能为null!");
        }
    }
}

```

```

    } else {
        this.cl_1 = cs.getCollection(CL_NAME_1);
    }
    if (null == cs) {
        throw new BaseException("集合空间不能为null! ");
    } else {
        this.cl_2 = cs.getCollection(CL_NAME_2);
    }
}

@Override
protected void reduce(Text key, Iterable<BSONWritable> values,
    Context context)
    throws IOException, InterruptedException{
    Iterator<BSONWritable> iterator=values.iterator
();
    long sum=0;
    List<BSONWritable> oldList = new ArrayList<BSONWri
table>();

    while(iterator.hasNext()){
        BSONWritable bsonWritable = iterator.next();
        oldList.add(bsonWritable);
    }

    //对values进行排序，排序字段为TRN_TIME（交易时间）
    Collections.sort(oldList, new Comparator<BSONWri
table>() {
        @Override
        public int compare(BSONWritable o1, BSONWritab
le o2) {
            String trn_time1 = (String)o1.getBson().get
("TRN_TIME");
            String trn_time2 = (String)o2.getBson().get
("TRN_TIME");
            return trn_time2.compareTo(trn_time1);
        }
    });

    Map<String,BSONWritable> result = new HashMap<St

```

```

ring, BSONWritable>());
    if (oldList != null && oldList.size() > 0) {
        //记录同一账户满足条件的笔数
        Map<String, BSONWritable> tempMap = new HashMap<String, BSONWritable>();
        for (int i=0; i<oldList.size()-1; i++) {
            BSONWritable bSONWritable1 = oldList.get
            (i);

            //交易代码
            String trn_cd = (String)bSONWritable1.getBson().get("TRN_CD");
            if (trn_cd.equals("000045") || trn_cd.equals("001045")
                || trn_cd.equals("021031") || trn_cd.equals("020031")
                || trn_cd.equals("001060") || trn_cd.equals("000060")) {

                //交易柜员
                String tran_teller_no1 = (String)bSONWritable1.getBson().get("TRAN_TELLER_NO");

                //流水号
                String jrn1_no = (String)bSONWritable1.getBson().get("JRN1_NO");

                //交易日期
                String trn_date1 = (String)bSONWritable1.getBson().get("TRN_DATE");

                //交易时间
                String trn_time1 = (String)bSONWritable1.getBson().get("TRN_TIME");
                Date bigDate = null;
                try {
                    bigDate = df.parse(trn_date1+" "+trn_time1);
                } catch (ParseException e) {
                    e.printStackTrace();
                }
                tempMap.put(jrn1_no, bSONWritable1);
                for (int j=i+1; j<oldList.size(); j++) {

```

```

        BSONWritable bsonWritable2 = oldList.
get(j);

        //交易代码
        String trn_cd1 = (String)bSONWritable
2.getBson().get("TRN_CD");
        if(trn_cd1.equals("000045") || trn_cd
1.equals("001045")
            || trn_cd1.equals("021031") ||
trn_cd1.equals("020031")
            || trn_cd1.equals("001060") ||
trn_cd1.equals("000060")){
            //交易柜员
            String tran_teller_no2 = (String)b
SONWritable2.getBson().get("TRAN_TELLER_NO");
            //流水号
            String jrn1_no2 = (String)bSONWrit
able2.getBson().get("JRN1_NO");
            //交易日期
            String trn_date2 = (String)bSONWri
table2.getBson().get("TRN_DATE");
            //交易时间
            String trn_time2 = (String)bSONWri
table2.getBson().get("TRN_TIME");
            Date smallDate = null;
            try {
                smallDate = df.parse(trn_date1+
" "+trn_time1);
            } catch (ParseException e) {
                e.printStackTrace();
            }
            //判断是否是同一实体{交易柜员}办理
            if(!tran_teller_no1.equals(tran_te
ller_no2)){
                continue;
            }
            //判断{交易日期}{交易时间}是否是[1小时]

```

内

```

        if((bigDate.getTime()-smallDate.ge
tTime())/1000 > tradeTime1){
            break;
        }
        tempMap.put(jrnl_no2,bSONWritable
2);

        }else{ //end if TRN_CD1.equals("00004
5")

            continue;
        }
    }//end for

    if(tempMap.size() >= tradeNum1){
        result.putAll(tempMap);
        tempMap.clear();
    }
    }else{
        continue;
    }//end if ||
} //end for
}

Map<String,BSONWritable> result2 = new HashMap<S
tring,BSONWritable>();

List<BSONObject> cl_1_list = new ArrayList<BSONO
bject>();

//结果写入sdb
Iterator iter1 = result.keySet().iterator();
while(iter1.hasNext()){
    String keyValue = (String)iter1.next();
    BSONWritable resultValue = result.get(keyValu
e);

    cl_1_list.add(resultValue.getBson());
    cl_1.insert(resultValue.getBson());
}
cl_1.bulkInsert(cl_1_list, DBCollection.FLG_INS
ERT_CONTONDUP);
cl_1_list = null;
List<BSONObject> cl_2_list = new ArrayList<BSONO
bject>();
context.write(null,null);

```

